

C++ ABI: the only thing that is more important than performance

iCSC 2020

Nis Meinert

German Aerospace Center



Deutsches Zentrum
für Luft- und Raumfahrt



Reading x86-64 Assembly

...for fun and profit

Function Prologue & Epilogue

- Few lines of code at the beginning (*prologue*) and end (*epilogue*) of a function, which **prepares** (and eventually restores)
 - the **stack** and
 - **registers**
- Not part of assembly: **convention** (defined & interpreted differently by different OS and compilers)

Prologue

```
1  push rbp      ; rbp: frame pointer
2  mov rbp, rsp ; rsp: stack pointer
3  sub rsp, N
```

alternatively

```
1  enter N, 0
```

(reserve N bytes on stack for local use)

Epilogue

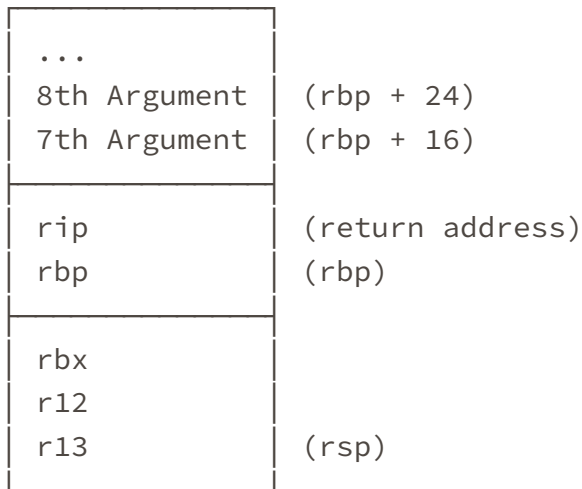
```
1  mov rsp, rbp
2  pop rbp
3  ret
```

alternatively

```
1  leave
2  ret
```

Stack frame for function call

- CALL = PUSH *address of next instruction* + JMP target
- RET pops return address and transfers control there
- pass arguments 1 ...6 in registers (rsi, rdx, ...)



(stack frame for function call with 8 arguments and local registers rbx, r12 and r13)

Reading assembly for fun and profit

```
1 int f(int x, int y, int z) {  
2     int sum = x + y + z;  
3     return sum;  
4 }
```

godbolt.org/z/MaWcP9

```
# g92 -O0  
| _Zlfiir:  
1| push rbp  
1| mov rbp, rsp  
1| mov DWORD PTR [rbp-20], edi  
1| mov DWORD PTR [rbp-24], esi  
1| mov DWORD PTR [rbp-28], edx  
2| mov edx, DWORD PTR [rbp-20]  
2| mov eax, DWORD PTR [rbp-24]  
2| add edx, eax  
2| mov eax, DWORD PTR [rbp-28]  
2| add eax, edx  
2| mov DWORD PTR [rbp-4], eax  
3| mov eax, DWORD PTR [rbp-4]  
4| pop rbp  
4| ret
```

godbolt.org/z/MaWcP9

Reading assembly for fun and profit

```
1 int f(int x) {  
2     return x + 1;  
3 }  
4  
5 int g(int x) {  
6     return f(x + 2);  
7 }
```

godbolt.org/z/87GK4q

```
# g92 -O0  
| _Z1fi:  
1|     push rbp  
1|     mov rbp, rsp  
1|     mov DWORD PTR [rbp-4], edi  
2|     mov eax, DWORD PTR [rbp-4]  
2|     add eax, 1  
3|     pop rbp  
3|     ret  
| _Z1gi:  
5|     push rbp  
5|     mov rbp, rsp  
5|     sub rsp, 8  
5|     mov DWORD PTR [rbp-4], edi  
6|     mov eax, DWORD PTR [rbp-4]  
6|     add eax, 2  
6|     mov edi, eax  
6|     call _Z1fi  
7|     leave  
7|     ret
```

godbolt.org/z/87GK4q

Reading assembly for fun and profit

```
1 void side_effect();  
2  
3 int f(int x) {  
4     auto a = x;  
5     side_effect();  
6     return a - x;  
7 }
```

godbolt.org/z/5xq5n5

```
# g92 -O0  
| _Z1fi:  
3| push rbp  
3| mov rbp, rsp  
3| sub rsp, 32  
3| mov DWORD PTR [rbp-20], edi  
4| mov eax, DWORD PTR [rbp-20]  
4| mov DWORD PTR [rbp-4], eax  
5| call _Z11side_effectv  
6| mov eax, DWORD PTR [rbp-4]  
6| sub eax, DWORD PTR [rbp-20]  
7| leave  
7| ret
```

godbolt.org/z/5xq5n5

Name mangling: C++ vs C

```
1 int f(int x) {  
2     return x * x;  
3 }  
4  
5 extern "C" int g(int x) {  
6     return x * x;  
7 }
```

godbolt.org/z/cj7bqx

```
# g92 -O0  
| _Z1fi:  
1 | push rbp  
1 | mov rbp, rsp  
1 | mov DWORD PTR [rbp-4], edi  
2 | mov eax, DWORD PTR [rbp-4]  
2 | imul eax, eax  
3 | pop rbp  
3 | ret  
|  
| g:  
5 | push rbp  
5 | mov rbp, rsp  
5 | mov DWORD PTR [rbp-4], edi  
6 | mov eax, DWORD PTR [rbp-4]  
6 | imul eax, eax  
7 | pop rbp  
7 | ret
```

godbolt.org/z/cj7bqx

Name mangling: C++ vs C

```
1 int f(int x) {  
2     return x * x;  
3 }  
4  
5 extern "C" int g(int x) {  
6     return x * x;  
7 }
```

godbolt.org/z/cj7bqx

Why?

- overloading
- namespaces
- templating

(Name of function doesn't suffice to resolve JMP location)

```
# g92 -O0  
| _Z1fi:  
1 | push rbp  
1 | mov rbp, rsp  
1 | mov DWORD PTR [rbp-4], edi  
2 | mov eax, DWORD PTR [rbp-4]  
2 | imul eax, eax  
3 | pop rbp  
3 | ret  
| g:  
5 | push rbp  
5 | mov rbp, rsp  
5 | mov DWORD PTR [rbp-4], edi  
6 | mov eax, DWORD PTR [rbp-4]  
6 | imul eax, eax  
7 | pop rbp  
7 | ret
```

godbolt.org/z/cj7bqx

Name mangling in C++

```
1 void f(int) {}  
2  
3 void f(double) {}  
4  
5 namespace my_fancy_namespace {  
6 void f(int) {}  
7 } // my_fancy_namespace
```

godbolt.org/z/jWY14x

```
# g92 -O2  
| _Z1fi:  
1| ret  
| _Z1fd:  
3| ret  
| _ZN18my_fancy_namespace1fEi:  
| ret
```

godbolt.org/z/jWY14x

- C++ does not standardize name mangling
- *Annotated C++ Reference Manual* even actively discourages usage of common mangling schemes. (Prevent linking when other aspects of ABI are incompatible.)

What is ABI?

What is ABI (**A**pplication **B**inary **I**nterface)?

Specifies interaction of functions and types across TUs[†] (translation units)

- Platform-specific (e.g., Linux)
- Vendor-specified (e.g., gcc)
- not controlled by WG21

Similar to a binary network protocol (Titus Winters)



Photo by Spencerian at
en.wikipedia.org (2005)

[†] *TU*: ultimate input to the compiler from which an object file is generated (*i.e.*, typically the `.cpp` file)

What is ABI (**A**pplication **B**inary **I**nterface)?

Specifies interaction of functions and types across TUs[†] (translation units) covering:

- Name mangling of functions
- Name mangling of types
- `sizeof` and alignment of objects
- Bytes semantics of the binary representation of objects
- Calling convention

Similar to a binary network protocol (Titus Winters)



Photo by Spencerian at
en.wikipedia.org (2005)

[†] *TU*: ultimate input to the compiler from which an object file is generated (*i.e.*, typically the `.cpp` file)

Why should I care?

...do you depend on any pre-compiled shared library?

Why should I care?

Why should I care?

- **Linking** different TUs requires usage of same ABI
- Typically a problem at API boundaries when combining TUs (e.g., shared libraries) that were compiled at different **times**
- Similar to binary network protocols: ABI tells you how to interpret bytes

Why should I care? \Leftrightarrow Why do network protocols have versions?

(Problem: ABI does not encode version number)

ABI does not encode version number

- **Q:** How to check if a given TU uses a compatible ABI?
- **A:** You can't!
- What happens if ABI is incompatible?
 - (a) Linking fails during compile time (good)
 - (b) Program spectacularly dies during run time (bad)
- Why isn't this a common problem?
 - Itanium ABI is mostly stable since C++11

History



ABI breakage of `std::string`

- Before C++11: `libstdc++` relied on copy-on-write (COW)
- C++11 disallows COW
 - fewer indirections
 - short string optimization (SSO)
- Problem: passing COW string to impl that expects SSO **may link** (same mangled name!)
 - one word passed
 - three words read
- *Solution*[†]: gcc changed mangled name

```
1 // until C++11
2 struct string {
3     struct control_block {
4         /* ... */
5     };
6     control_block *data;
7 };
8
9 // since C++11
10 struct string {
11     char *data;
12     std::size_t size;
13     std::size_t capacity;
14 }
```

godbolt.org/z/KM5Tvq

↪ Take-away for compiler vendors: ABI break was a huge disaster

[†] RHEL 7 still uses old `std::string` ABI to provide compatibility for older .so

Quiz time